

Week 10, topics:

- Monte Carlo for Statistical Mechanics
 - Markov Chain
 - Simulated Annealing
 - Ising Model example

Monte Carlo simulations:

Any simulation that uses random numbers to simulate random physical processes to estimate something about the outcome of that process.

Newman, p. 476.

We focus on statistical mechanics here.

Statistical mechanics: a review

- For a system in equilibrium at temperature T , the probability of finding the system in any particular state i is given by the Boltzmann distribution,

$$P(E_i) = \frac{\exp[-E_i/(k_B T)]}{Z}, \quad Z = \sum_{i=1}^{ALL} \exp[-E_i/(k_B T)]$$

where E_i is the energy of state i , k_B is Boltzmann's constant, Z is partition function

- System at temperature T undergoes transitions between states, with probability of being in a particular state $P(E_i)$

- To calculate expectation value of a macroscopic property X (e.g. total energy, magnetization, spin): average over the many microstates that the system visits

$$\langle X \rangle = \sum_{i=1}^{ALL} X_i P(E_i)$$

where X_i is the value of the quantity in the i^{th} microstate and P is the probability of finding the system in that microstate.

- Simple example: single mole of gas has $N_A \approx 6 \times 10^{23}$ molecules. Assume each molecule had only 2 possible spin states (gross underestimation), then the total number of spin microstates of the mole of gas is 2^{N_A} , which is huge.
- If we want to get any useful macroscopic information $\langle X \rangle$ about the system, we need to be more clever about how we count than just brute-force counting everything.

$$\langle X \rangle = \sum_{i=1}^{ALL} X_i P(E_i), \quad P(E_i) = \frac{\exp[-E_i/(k_B T)]}{Z}, \quad Z = \sum_{i=1}^{ALL} \exp[-E_i/(k_B T)]$$

- Huge number of terms in sum \Rightarrow often impossible to perform analytically \Rightarrow use Monte Carlo summation.
- Two difficulties to overcome:
 1. estimating $\langle X \rangle$: properly sampling which terms to sum over (*solution: importance sampling*),
 2. estimating Z (*solution: Markov Chain Monte Carlo*)

Importance sampling for Stat. Mech.

- Randomly sample the terms in the sum and only use those as an estimate. Replace $\langle X \rangle = \sum_{i=1}^{ALL} X_i P(E_i)$ with a sum over N randomly sampled microstates,

$$\langle X \rangle = \frac{\sum_{i=1}^N X_i P(E_i)}{\sum_{i=1}^N P(E_i)}.$$

- Denominator ensures total probability over the sampled states is 1.
- To get a good estimate for the sum, it is only worth keeping the big terms
 - A lot of states have $E_i \gg k_B T$, therefore $P(E_i)$ really small since probability is proportional to $\exp[-E_i/(k_B T)]$
- Need to preferentially choose terms where the integrand is non-negligible, but assign them less weight individually (to not bias the final estimate) ... so use importance sampling!

- For a discrete sum (see p. 478),

$$\langle X \rangle = \sum_{i=1}^N X_i P(E_i) \approx \frac{1}{N} \sum_{k=1}^N \frac{X_k P(E_k)}{w_k} \sum_{i=1}^{ALL} w_i.$$

- For weight w , use $P(E_i)$

$$\langle X \rangle \approx \frac{1}{N} \sum_{k=1}^N \underbrace{\frac{X_k P(E_k)}{P(E_k)}}_{=X_k} \underbrace{\sum_{i=1}^{ALL} P(E_i)}_{=1}.$$

In the end,

$$\langle X \rangle \approx \frac{1}{N} \sum_{k=1}^N X_k.$$

- Looks simple and no different from mean-value MC,
- but recall that the X_k 's are drawn from non-uniform distribution: we randomly choose terms in the sum based on their Boltzmann probabilities, $P(E)$. But how? Recall

$$P(E_i) = \frac{\exp[-E_i/(k_B T)]}{Z}, \quad Z = \sum_{i=1}^{ALL} \exp[-E_i/(k_B T)]$$

- To do it this way, we need Z , which is a sum over all states. But if we could do this, we wouldn't need Monte Carlo in the first place!

Markov chain method

Mish-mashing https://en.wikipedia.org/wiki/Markov_chain (https://en.wikipedia.org/wiki/Markov_chain) and https://en.wikipedia.org/wiki/Markov_property (https://en.wikipedia.org/wiki/Markov_property).

A Markov chain is a stochastic model describing a sequence of possible events in which the probability of each event depends only on the state attained in the previous event. [...] (sometimes characterized as "memorylessness"). In simpler terms, it is a process for which predictions can be made regarding future outcomes based solely on its present state [...]. In other words, conditional on the present state of the system, its future and past states are independent.

- Random walks (Brownian motion) are Markov chains.
- Here: events are jumps in energy states, one after another.

Solution: Use the Markov chain method.

- Text goes into details on how to implement this method with a Metropolis algorithm.
- Crucial key: Metropolis does not directly compute probability to be in one state, but instead uses probability to transition between two states
 - Z cancels out in the process!
- In this lecture, we will summarize it algorithmically first, then briefly outline why it works mathematically.

Algorithm

1. Choose a random starting state i
2. Calculate the energy of that state E_i
3. Choose a transition to a new state j uniformly at random from allowed set
4. Calculate the energy of this new state, E_j
5. Calculate the acceptance probability for this transition:
 - $P_a = 1$ if $E_j \leq E_i$ (always accept a lower energy state)
 - $P_a = \exp\left(-\frac{E_j - E_i}{k_B T}\right)$ if $E_j > E_i$ (accept a higher energy state sometimes, more often for high T).
6. Accept/reject the move according to the acceptance probability
7. Measure the quantity X you want in its current state (new or old i) & store it
8. Repeat from step 2.

- How to implement the probability of the event in the previous slide?
 - $P_a = 1$ if $E_j \leq E_i$ (always accept a lower energy state)
 - $P_a = \exp\left(-\frac{E_j - E_i}{k_B T}\right)$ if $E_j > E_i$ (sometimes accept a higher energy state, more often for high T).
- Draw a random number in $[0, 1)$. The statement


```
if random() < exp(-(Ej-Ei)/kT):
```

 will introduce what to do if the move is accepted (`elif` will introduce what to do if rejected).
 - E.g., at very high T , $\exp(\dots) \approx 1$ and almost all moves are accepted.
 - E.g., at low T , say, $\exp[-\Delta E/(k_B T)] = 1\%$, then `random()` has 1% chance of drawing a number that is $< 1\%$.
 - If $E_j \leq E_i$, then $\exp(\dots) \geq 1$ and `if` statement automatically accepts.

Why does Metropolis create a system where each microstate has a probability $P(E_i)$, the Boltzmann distribution?

- Let τ_{ij} be transition probability from microstate i to microstate j , such that

$$\frac{\tau_{ij}}{\tau_{ji}} = \frac{P(E_j)}{P(E_i)},$$
 which the Metropolis algorithm satisfies (see p. 482). For example, a small ratio above means
 - RHS: " j is much less probable than i " and equivalently,
 - LHS: " $\text{probability of } j \rightarrow i \text{ is much higher than } i \rightarrow j$ ",
 both statements being true in equal amounts.
- "In equal amounts" is crucial: it means that if you start from any initial state, your system will progressively evolve towards one where all microstates follow Boltzmann.

- Does it always converge? Yes. If you love linear algebra, see proof in Appendix D of Newman (it is actually quite elegant).

Example: Ising model


- Simple model of ferromagnetism, demonstrating many physical characteristics of fancier models.
- Assume an object is made up of a collection of dipoles (e.g., electron spins) and the net magnetization is the sum of the magnetization of all the spins.
- Ising model:
 - assume the spins can only point up or down.
 - the spins interact and favour parallel alignment of pairs of spins
 - the interactions are non-zero only between nearest neighbours (i.e., distance dependent).
- Macroscopic energy E and magnetization M (with no external field) are given by

$$E = -J \sum_{\langle ij \rangle} s_i s_j \quad \text{and} \quad M = \sum_i s_i$$

where $s = +1$ if spin is up & $s = -1$ if spin is down.

- Lowest energy occurs if the spins all line up.

- Spins can randomly flip as the system visits a set of allowable states given its temperature. At any particular moment the system may look like



Example in 1D

- Create array of dipoles, initial state: random spin at each location.
- Calculate energy & magnetization of state
- Implement Metropolis algorithm:
 - create new state: flip 1 spin randomly
 - calculate new total energy
 - calculate acceptance probability
 - decide whether to accept or reject new state
 - store 'new' energy & magnetization
 - repeat

```
In [ ]: """ This program calculates the total energy and magnetization
for a 1D Ising model with N dipoles
Author: Nico Grisouard, University of Toronto """

import numpy as np
from random import random, randrange
import matplotlib.pyplot as plt

def energyfunction(J_, dipoles):
    energy = -J_*np.sum(dipoles[0:-1]*dipoles[1:])
    return energy

def acceptance(ARGUMENTS):
    """ Function for acceptance probability; to be completed """
    # Do stuff here
    return result # result is True or False
```

```
In [ ]: # define constants
kB = 1.0 # "Boltzmann" constant
T = 1.0 # "Temperature"
J = 1.0 # Interaction
num_dipoles = 100
N = 100 # number of flips

# generate array of dipoles and initialize diagnostic quantities
dipoles = np.ones(num_dipoles, int) # hint: this will not work
energy = []
magnet = []
E = energyfunction(J, dipoles)
energy.append(E)
magnet.append(sum(dipoles))
```

```
In [ ]: print(dipoles)
```

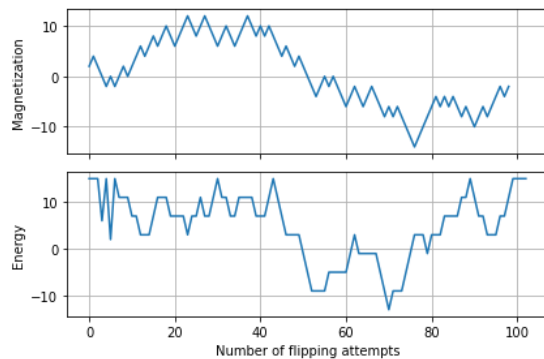
```
In [ ]: for i in range(N):
    picked = randrange(num_dipoles) # choose a dipole
    dipoles[picked] *= -1 # propose to flip this dipole
    Enew = energyfunction(J, dipoles) # compute Energy of proposed new state

    # calculate acceptance probability
    accepted = acceptance(Enew, E,)

    # store energy and magnetization
```

```
In [15]: # plot energy, magnetisation
fg, ax = plt.subplots(2, 1, sharex=True)
ax[0].plot(magnet)
ax[0].set_ylabel('Magnetization')
ax[0].grid()
ax[1].plot(energy)
ax[1].set_xlabel('Number of flipping attempts')
ax[1].set_ylabel('Energy')
ax[1].grid()


plt.tight_layout()
plt.show()
```



Simulated annealing

- Using Monte Carlo simulations to find **global** minima/maxima.
- In week 4 we talked about ways of finding local minima, which worked well if we had one, well-bracketed minimum. How about multiple local minima?
- How it works: rewrite max/min problem as looking for a "ground state energy" of a system.
 - Function f that you want the max/min of: make this the energy function.
 - how could you find ground state: reduce temperature until you reach the ground state.
- Issue: if you reduce temperature too quickly: might get caught in a local min instead of the global min.
- Solution: reduce temperature slowly. This way system has time to explore many microstates and find a good approximation to the global minimum.


- Visual Analogy: particle in a bumpy potential.
 - Too low energy: get stuck in nearest local minimum.
 - Keep low energy but allow some random 'kicks' in energy: can kick out of local minimum and continue heading to global minimum
 -



(https://commons.wikimedia.org/wiki/File:Hill_Climbing_with_Simulated_Annealing.gif)
(https://commons.wikimedia.org/wiki/File:Hill_Climbing_with_Simulated_Annealing.gif)

Example: travelling salesman

- Famous NP-hard problem (<https://en.wikipedia.org/wiki/NP-hardness> (<https://en.wikipedia.org/wiki/NP-hardness>)): what is the shortest route to visit a given set of locations on a map?
- Want global minimum of distance
- Start with random route, swap 2 cities, use Metropolis algorithm to determine whether to keep the swap
 - "energy" in this case is the total distance of the route
 - You can explore this problem using code from the book (`salesman.py`).



Summary

- Stat mech results often involve sums over impractical number of terms
 - Monte Carlo can sample a representative number of microstates: importance sampling
 - weight = probability of given microstate
 - How to compute partition function: Markov Chain using Metropolis algorithm.
 - Randomly explore states by transitioning from one microstate to the next
 - Works because the transition probability is closely related to the ratios of probabilities of being in either state.
 - Famous and fundamental example: the Ising model.
- Simulated annealing for finding global minimum of a function: let f be the "energy", and lower the "temperature" progressively
 - Famous example: travelling salesman problem